

Multi-modal fusion for violent scene detection in Hollywood movies

Elsevier¹

Radarweg 29, Amsterdam

Elsevier Inc^{a,b}, Global Customer Service^{b,}*

^a1600 John F Kennedy Boulevard, Philadelphia

^b360 Park Avenue South, New York

Abstract

Violent scene detection (VSD) is a challenging problem because of the heterogeneous content, large variations in video quality, high semantic meaning of the concepts. In the last few years, combining multiple features from multi-modalities is proven as an effective strategy for a wide range of video classification tasks and almost existing works have focused on general multimedia event detection (MED), the specific event detection like VSD has been comparatively less studied. In this work, we try to study the value of multi-modal features and combination of these features for VSD in Hollywood movies. We rigorously analyze and combine a set of low-level features and deep learning feature that capture appearance, color, texture, motion and audio in videos. We also evaluate the utility of mid-level visual information obtained from detecting related violent concepts. Experiments are performed on MediaEval VSD 2014 dataset, made publicly available. Results show the performance of visual and motion features are better than audio features. The performance of mid-level features was nearly as good as that of low-level visual features. Experiments with a number of fusion methods show that all single features are complementary and help to

[☆]Fully documented templates are available in the elsarticle package on CTAN.

^{*}Corresponding author

Email address: support@elsevier.com (Global Customer Service)

URL: www.elsevier.com (Elsevier Inc)

¹Since 1880.

improve overall performance. This study also provides an empirical foundation for selecting feature sets that are capable of dealing with heterogeneous content data as violent scenes in movies.

Keywords: violent scene detection, video retrieval, multi-modal fusion, multiple features

1. Introduction

Nowadays, the movie industry generates thousands of movies each year. However, not all movies are suitable for young people (especially children, teenagers, e.g.) to watch, because they might have violent contents. The re-
search from University of Pittsburgh ² reports that watching violence in movies
or TV programs tends to make children more aggressive and leads to unhealthy
attitudes. It is crucial to have violent scene detection (VSD) system which en-
ables the parents to choose movies that are suitable for their children. More
general, for content provider, the violent scene detection technique can be used
to assist in movie rating; for general end users, it can block the violent content
in client terminal devices.

Violent scene detection in the movie is a very challenging problem, particularly because the movies contain the heterogeneous content, have large variations in quality. Because violence concept has a high semantic meaning, violent

²<http://www.ocd.pitt.edu/Files/PDF/Parenting/TvAndMovieViolence.pdf>

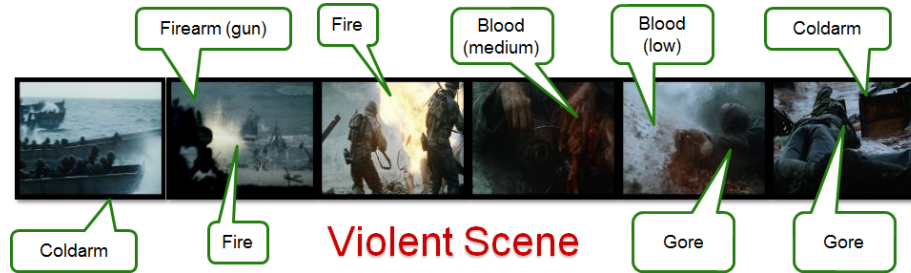


Figure 1: Example of a Violent Scene in Saving Private Ryan Movie.

15 segments in video can contain the huge collection of **affective content** which refers to the characterization of violence concept, such as blood, gun, cold arm, firearm; **actions** such as fighting, running, shooting, etc.; **activities** such as car chase, boat chase, etc.; **scenes** such as gory scene, overcast, scene etc.; **sound** such as gun shot, scream, explosion, etc.; and their strictly spatial/temporal
 20 relationships (for example: a scene with two man are fighting first and after that they get injury or pain). Besides that, movies often have a lot of effects, scene transition techniques which are heavily edited by the producers. Figure 1 shown an example of violent scene in Hollywood movie.

The goal of this paper is to develop a VSD system as shown in Figure 2. In
 25 this system, movie will be divided into shots and then shots are classified and scored by using multi-modal information. A general framework to build a such system usually consists of 3 components: feature extraction, learning model, and prediction. Most of existing approaches focus on feature extraction component and many features such as HOG, HOF, SIFT, GIST, LBP are evaluated in
 30 the frameworks such as multimedia event detection, action recognition, scene classification. The main observation drawn from these studies is how to select the best features and combine of multi modal features to handle large variations of violent scenes.

In this paper, we present a framework that is designed to evaluate every

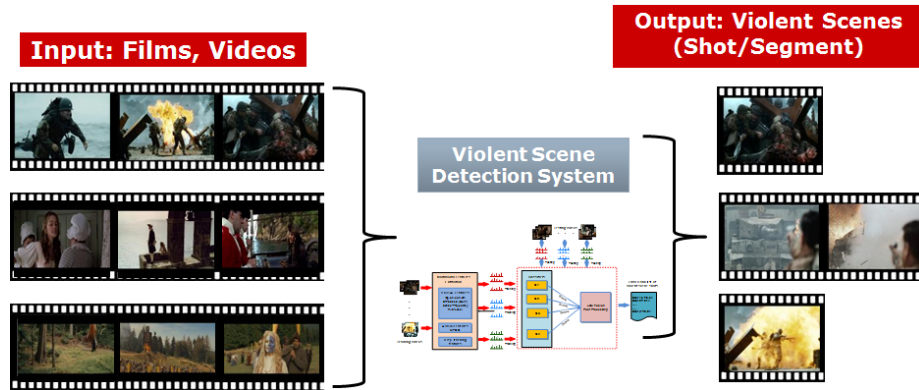


Figure 2: Overview of violent scene detection system.

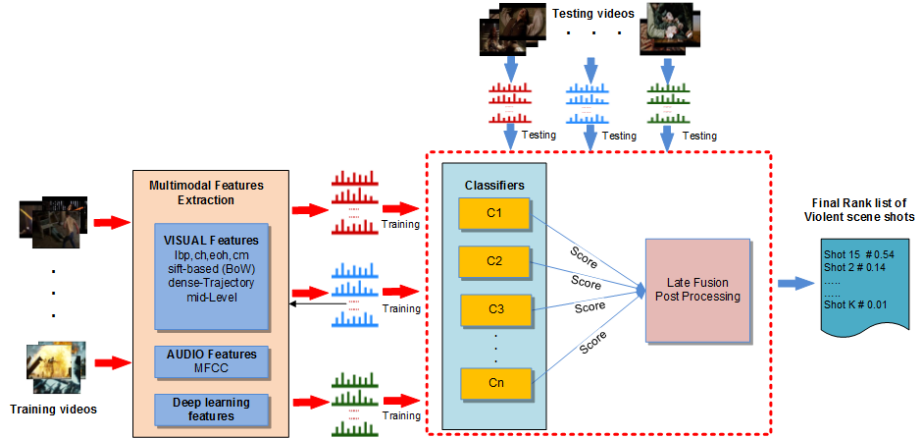


Figure 3: Overview of violent scene detection system.

single type of features and incorporate multi-modal features for VSD. Figure 3 shows our system architecture. We extracted visual (image and video), audio and deep learning features. Visual features are proposed to describe visual characteristics of violent concept, such as the objects, the scenes. Audio feature is used to capture the sound of violent actions. In addition, motion feature is used describe the actions, the activities. We rigorously analyze and combine a set of low-level features that capture violent appearance in video shots. We also evaluate the utility of mid-level representation to reduce the semantic gap between violence concept and low-level features. Mid-level means this layer lies between low-level features and high-level target events. Further, we investigate how to combine these diverse features using several fusion methods to improve overall performance of system. The main focus of this work is to provide a comparative analysis of value of multi-modal features for VSD. Our study is helpful in that it evidences strong and weak points of current violence detection techniques and can be used to guide further work in the area.

We perform our experiments of a standard benchmark, MediaEval VSD 2014[1] which are made publicly available³. Our system is trained and tested

³<http://www.technicolor.com/en/innovation/research-innovation/scientific-data->

on Hollywood movies within MediaEval VSD Dataset and evaluated by MediaEval VSD official metric (MAP-mean average precision). By using this standard dataset and evaluation protocol makes our proposed methods to be fair to compare with results by other MediaEval VSD participants (results did not use the external data). Results show the performance of visual features are the best. Low-level visual features and motion feature play very important role when fusing all modalities. Audio and mid-level features are not good as low-level visual features but they are complementary to improve overall performance. Experiments with several fusion methods show that all single features are complementary and help to improve overall performance.

The rest of this paper is organized as follows. We review some related work in section 2. Section 3 introduces our framework, our approaches for feature extraction, and detail of feature configurations, fusion methods. The experimental results and their analysis will be in section 4, as well as the conclusion and future works are described in the section 5 of this paper.

2. Related work

Violent scene detection is a kind of multimedia event detection. Although combining multiple features from multi-modalities has been proven as an effective strategy for multimedia event detection, relatively few works have been proposed to apply these approaches for violent scene detection in video. The main reason is that the definition of violence is ambiguous. It is difficult to describe this high-level concept using mathematical formulation precisely. In general, the violence concept is not well-defined and recent approaches addressed the problem by their own definitions.

Some of the previous works applied different kind of visual features to detect flame, blood, explosion as the informative cues for violence. As a first approach in this field, Jeho [2] propose an approach to recognize violent scenes

sharing/violent-scenes-dataset

in videos by detecting flame and blood, capturing the degree of motion activity,
80 the soundtrack and characterization of sound effects of violent events. Meanwhile, Chen [3] decompose violent scene detection into action scene detection and bloody frame detection. Clarin [4] present a system which uses a Kohonen self-organizing map, which is for detecting skin and blood pixels in each frame, and motion intensity analysis to detect violent actions involving blood. More
85 recently, Gong Yu [5] introduce a violence detector using low-level visual and auditory features and high-level audio effects to identify potential violent content in movies. Jian [6] describe a weakly-supervised audio violence classifier combined using co-training with a motion, explosion and blood video classifier to detect violent scenes in movies. Penet [7] compared two modality fusion
90 methods, namely Early Fusion and Late Fusion. Early Fusion concatenates features from both modalities before machine learning, while Late Fusion fuses probabilities of both modalities already calculated. They reported Late Fusion was superior to Early Fusion.

Beside low-level based approaches, violent scene detection is also a kind of
95 high level recognition task. The major challenge is to deal with the gap of semantic meaning. In recent years, there are some approaches using attributes to narrow the semantic gap in high-level recognition task, such as object recognition using attributes, scene classification using attributes, action recognition using actions bank[8]. Li [9] propose Object Bank (OB), a new representa-
100 tion of natural images based on objects (using objects as attributes to describe scenes). Liu [10] use high-level semantic concepts, also called attributes, to represent human actions from videos and argue that attributes enable the construction of more descriptive models for human action recognition. Sasanand [8] present Action Bank, a new high-level representation (kind of attributes) of
105 video. For VSD, Bodgan [11] rely on fusing mid-level concept predictions made using multi-layer perception classifiers to automatically localize the occurrence of violence within a video. They proposed a frame-level violence prediction, applying a multi-layer perceptron in order to utilize these concepts. They put the first layer for the concept prediction, and the second layer for the violence

110 prediction. In addition to those provided concepts, Tan [12] have utilized extra
42 violence concepts such as bomb and war from ConceptNet[13]. ConceptNet
is composed of nodes representing concepts in the form of words or short phrases
with their relationships. On their system those extra concepts are trained using
YouTube videos which are crawled additionally.

115 In the last few years, MediaEval Violent Scene Detection affect Task is be-
come popular, many state-of-the art systems have been developed and reported
recently[1]. Most systems extract a large number of multimodal features from
visual, audio signals and text. The features include low-level features such as
color histogram, edge of histogram, local binary pattern, SIFT based, Space-
120 Time Interest Points (STIP), Histograms of Oriented Gradients (HoG), His-
tograms of Optical Flow (HoF), Motion Boundary Histograms (MBH), visual
activity, MFCCs, trajectory-based features (cite each team). Different state-of-
the-art features encoding are used, such as Bag-of-Visual-Words representation,
Fisher vector representations. In addition, some approaches use of pre-defined
125 concept detectors like part-level attributes FUDAN Team[14], mid-level con-
cept description, with provided concepts from ConceptNet (e.g., punishment,
victim, rape, etc) VIREO team[12]. The video shots are classified by using dif-
ferent machine learning techniques (Support Vector Machines (SVMs), k nearest
neighbors (kNN), Bayesian Network). Multimodal integration is achieved via
130 early fusion[15] and late fusion[15, 16, 17, 14] where score fusion is often used
to combine scores independently computed from different subsets of features.

In this paper, we present a novel violent scene detection framework which
is designed to incorporate many of the above-mentioned success principles in
its architecture. In particular, our work incorporates novel developments into
135 the system, which can be summarized into major contributions. First, our
work developed and incorporated multi modal features at diverse granularities
to evaluate the value of different types of features for VSD. Second, our work
explores the use of mid-level concept features, which are detected based on
low-level features, aiming to provide more semantic understanding capability
140 into our system. Third, we present several approaches to learn fusion functions

which combine violent scores (late fusion) to improve overall performance.

3. VSD system

3.1. Framework Overview

We use a unified framework to evaluate the performance of each feature
145 as well as the performance of feature combination methods. The framework
should be flexible so that we can easily test different kind of features or fusion
strategies. It also should be designed in a component-based manner so that we
can evaluate each component separately while keeping other components intact.
In this spirit, our framework can be decomposed into following components:
150 (1) Pre-processing, (2) Feature Extraction, (3) Feature Encoding, (4) Feature
Classification and (5) Feature Fusion. The detail of each component is described
in the following sections. We especially focus on the feature encoding and feature
fusion component. The former is used to evaluate features while the later is used
to evaluate fusion strategies for VSD. An overview of our framework is shown
155 in Figure 3.

3.2. Pre-processing

To create a feature representation for each shot: firstly, we extract five
keyframes per second; secondly, we extract visual feature for each keyframe;
lastly, we apply maximum and average pooling methods to create shot-based
160 features.

3.3. Feature Extraction

The feature extraction component aims to make a discriminative vector rep-
resentation for each shot that is extracted from the pre-processing step. The
extraction method depends on what type of feature will be used. To conduct
165 a comprehensive evaluation of features for VSD, we support a large variety of
features including global and local visual features. Global features capture the
global statistics of each extracted shot. These statistics can be calculated di-
rectly from sub regions of a sampled frame and then concatenated to form the

vector representation for that frame, before being aggregated to make the final
 170 representation for each shot. It is more complicated to calculate the feature
 vector representation for local features. The number of local features vary from
 frame to frame, therefore it requires a special encoding technique which will be
 described in 3.4.

Besides global and local features, other features are also supported in our
 175 evaluation framework. Audio feature can be extracted from pre-defined tempo-
 ral windows. Feature of each temporal window provide a local audio character-
 istic at that temporal location. Therefore, audio feature can be considered as a
 local feature and can be well-integrated into our feature extraction framework.
 Another feature that is also supported is mid-level feature, which is represented
 180 from violent concept detectors. We consider both the in-domain and out-domain
 concept detectors. In case the concepts are used from off-the-shelf datasets, we
 employ the state-of-the-art deep learning features which are extracted from a
 pre-trained model. The detail description of each feature is presented in Section
 4.

185 3.4. Feature Encoding

We employed the bag-of-words model with codebook size of 1000 and soft-
 assignment technique to generate a fixed-dimension feature representation for
 each keyframe. Besides encoding the whole image, we also divided it into grids
 of 3x1 and 2x2 to encode spatial information. Besides the bag-of-words repre-




Figure 4: Bag of visual words.

190 sentation for the whole image, we divided each key frame into 2×2 and 1×3
 sub-regions to cover information about the spatial layout of features, which is

a trade off between performance and computation cost brought by the high dimensionality of the feature vector, and computed the bag-of-word representation for each sub-region. We applied this method for visual local features.

195 3.5. Feature Classification

LibSVM [18] is used for training and testing at shot level. To generate training data, shots which fall into positive segments more than 80% will be considered as positive shots. The remaining shots are considered as negative. Extracted features are scaled to $[0, 1]$ using the svm-scale tool of LibSVM.
200 We use chi-square kernel to calculate the distance matrix. The optimal (C;g) parameters for learning SVM classifiers are found by conducting a grid search with 5-fold cross validation on the original dataset.

3.6. Feature Fusion

3.6.1. Late Fusion

205 Fusing information coming from different media seems a natural way to handle multimedia content. Fusing of multi-modal information has been widely used for tasks including multimedia event detection, video search, etc. Naturally different semantic events, types of multimedia data have their own characteristics so their fusion strategies could be different also. For violent scene detection,
210 we investigated several simple but effective fusion schemes to construct a better violence prediction system that exploits the advantages of each type of modalities.

4. VSD features

4.1. Visual global features

215 We evaluate both global features and local features. We selected the best configuration of each types of features from our previous work[19]. The global features include color moments, color histogram, edge orientation histogram, and local binary patterns with different configuration:

- Granularity: Since global features do not capture spatial information,
220 to overcome this problem, a grid nm usually used to divide the input
image into non overlapping sub-regions. The features extracted from these
regions are concatenated to form the feature vector for the image.
- Color space: Local binary patterns (LBP) and edge orientation histogram
(EOH) are extracted from gray scale image. For color moments and color
225 histogram, color spaces including HSV, RGB, Luv, and YCrCb are used.
- Quantization: For color histogram, we only use 8-bin histogram for each
channel. For edge orientation histogram, we quantize orientations into
histograms of 36 (edge) + 1(non-edge) bins. For local binary patterns, we
quantize binary patterns into histograms of 30, 59 bins.

230 4.2. Visual local features

4.2.1. Still image feature

For local feature, we use popular SIFT with both Hessian Laplace inter-
est points and dense sampling at multiple scales. For dense sampling, besides
normal SIFT descriptor, we also use Opponent-SIFT and C-SIFT. For Harlap
235 interest point detector, we only use normal SIFT descriptor.

4.2.2. Motion feature

As shown by Wang et al. [20], Dense Trajectory feature is one of the best
approaches for action classification. Dense Trajectory has an efficient solution
to remove camera motion. In the Hollywood movies, there are a lot of action
240 with different movie effects in the violent scenes. We try to apply the dense
trajectory to capture these information. Trajectories are obtained by tracking
the densely sampled points in the optical flow fields. We use Motion Boundary
Histogram (MBH) to describe each trajectory. This feature descriptor is good
for handling camera motion.

245 4.2.3. Audio feature

We use the popular MFCC for extracting audio feature. We choose the audio segment of 25ms and step size of 10ms. The 13d MFCCs along with each first and second derivatives are used for representing each audio segment. Raw MFCC features are also encoded using BoW.

250 4.3. Mid-level feature

4.3.1. In domain: VSD concepts

Violent scene detection is also a kind of highlevel recognition task. Violence concept has high semantic meaning and high variability in appearance. Besides that, due to lack of training data for this concept, training violence classifier
255 directly from low-level features is not effective. Instead of focusing on selecting the low-level features for VSD, we try to investigate how to use related violent information as mid-level feature to detect violent scene in the movie. We manually select the violent information which are annotated by human assessors as related attributes of violence concept. Then we propose to use such concepts
260 as related semantic attributes of the original violence concept. Figure 1 shows how to describe the violent scene by their semantic attributes. These attributes which can be more well-presented by low-level features (e.g. the attribute about blood can be sufficiently presented by color) are used to create the mid-level features. In other words, the attributes have smaller gap to the low-level features, compared to the original violence concept. For example, in Figure 1, because
265 violence occurred during the whole scene, it is very hard to use low-level features to represent the violence concept. But, we can use the low-level visual features to represent some attributes such as fire, blood. By doing this, we narrow the semantic gap between the original violence concept and low-level features extracted from video shots. Attributes related to violence concept are manually
270 defined. We used seven related concepts provided in MediaEval Dataset as violent attributes. Our mid-level framework is shown in Figure 5. Firstly, each corresponding attribute classifier is trained by using low-level features. Secondly, the mid-level features of a training (or test) video shots are formulated

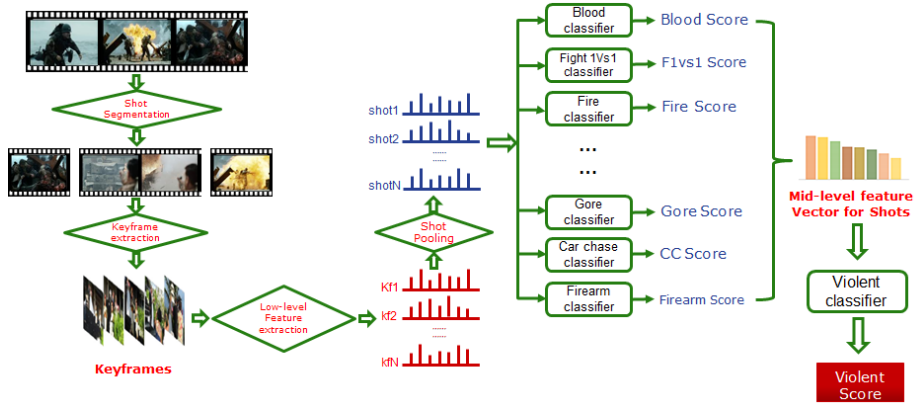


Figure 5: Overview of Mid-level framework.

by concatenating scores returned by attribute classifiers. At the moment, we use the same weights for concatenating different scores. Then thirdly, we use this feature to train the mid-level feature-based violence classifier. Finally, we apply this violence classifier on test set to get the violence score for each shot (these shots are also represented by mid-level features).

4.3.2. Out domain: Deep learning feature

5. Experiments

5.1. Dataset

In this paper, we used the dataset from MediaEval Affect Task 2014 [1], this is a set of 31 Hollywood movies that must be purchased the original DVD due to copyright issues. The movies are of different genres (from extremely violence movies to movies without violence). In this dataset, we focus on the violent concept with subjective definition, which is defined as those which one would not let an 8 years old child see because they contain physical violence. Follow the proposed method, we divide this dataset into 2 parts:

- DEVEL (Table 1): this is training dataset; used to train violence classifiers; has 24 movies, total 34779 shots, 48.29 hours.

No.	Video Name	Length (in second)	# keyframes	# shots
1	Armageddon-1998-dvd2002	8681.05	217026	1737
2	BillyElliot-2000-dvd2003	6349.36	158734	1270
3	Eragon-2006-dvd2007	5985.57	149639	1198
4	HarryPotter5-2007-dvd2008	7954.72	198868	1591
5	IAmLegend-2007-dvd2010	5780.58	144514	1157
6	Leon-1994-dvd2004	6344.49	158612	1269
7	MidnightExpress-1978-dvd2008	6960.96	174024	1393
8	PiratesOfTheCaribbean1-2003-dvd2006	8241.01	206025	1649
9	ReservoirDogs-1992-dvd2004	5712.98	142825	1143
10	SavingPrivateRyan-1998-dvd2006	9750.89	243772	1951
11	TheSixthSense-1999-dvd2000	6178.01	154450	1236
12	TheWickerMan-2006-dvd2008	5870.89	146772	1175
13	TheBourneIdentity-2002-dvd2006	6816.29	170407	1364
14	TheWizardofOz-1939-dvd2000	5859.29	146482	1172
15	DeadPoetsSociety-1989-dvd2002	7415.17	185379	1484
16	FightClub-1999-dvd2001	8006.34	200158	1602
17	IndependenceDay-1996-dvd2010	8834.96	220874	1767
18	TheGodFather-1972-dvd2008	10194.96	254874	2039
19	PulpFiction-1994-dvd2009	8887.97	222199	1778
20	ForrestGump-1994-dvd2006	8176.97	204424	1636
21	Fargo-1996-dvd2004	5646.34	141158	1130
22	ThePianist-2002-dvd2007	8567.1	214177	1714
23	FantasticFour1-2005-dvd2005	6094.41	152360	1219
24	LegallyBlond-2001-dvd2002	5523.49	138087	1105
	Total	173833.8	4345840	34779

Table 1: DEVEL set includes 24 Hollywood movies.

No.	Video Name	Length (in second)	# keyframes	# shots
1	V_FOR_VENDETTA	7626.49	190662	1526
2	TERMINATOR_2	8831.37	220784	1767
3	JUMANJI_COLLECTORS_EDITION	5993.98	149849	1199
4	GHOST_IN_THE_SHELL	4966	124150	994
5	DESPERADO	6012.89	150322	1203
6	BRAVEHEART	10224.49	255612	2045
7	8_MILE	6355.53	158888	1272
	Total	50010.75	1250267	10006

Table 2: TEST set includes 7 Hollywood movies.

- TEST (Table 2): this is testing dataset; used to test and evaluate the system; has 7 movies, total 10006 shots and 13.89 hours.

Total duration are about 62.18 hours, with 44.785 shots. To reduce the computation cost, when we extract keyframes, we resize the keyframes to 500x400 pixels.

5.2. Groundtruth

By using subjective definition in MediaEval VSD 2014[1], the ground truth is created by human assessors and provided by the MediaEval organizers. In addition to segments containing physical violence, annotations also include the following high-level concepts: presence of blood, fights, presence of fire, presence of guns, presence of cold arms, car chases and gory scenes, for the visual modality; gunshot, explosion and scream for the audio modality. The ground truth data are provides in segment. To generate training data, we consider the positive shots are the ones which have 80% overlapping with ground truth segments.

5.3. Evaluation Metrics

We use mean average precision (mAP), an evaluation metric that is widely used for classification and retrieval systems. The mAP is computed based on the ranked list of shots returned by the detection system and the ground truth provided by the task organizers. Mean average precision (mAP) is calculated as below:

$$MAP = \frac{\sum_{v=1}^V AP(v)}{V}$$

, where V is number of test videos and AP is average precision for each video.

5.4. Results and comparison

5.4.1. Evaluation of global features

5.4.2. Evaluation of local features

5.4.3. Evaluation of motion features

5.4.4. Evaluation of audio features

5.4.5. Evaluation of mid-level features

5.4.6. Evaluation of deep learning features

5.4.7. Evaluation of fusion schema

5.4.8. Comparison with MediaEval teams

6. Conclusion

We evaluate the performance of multi-modal features for the violent scene detection. The performance of global features and local features that are widely used in state of the art classification systems are compared. Our study can be served as a baseline for comparison of advanced algorithms or systems in the violent scene detection. Still image based feature (SIFT) has better performance than motion based feature (Dens-trajectory + MBH). Fusion of global features is effective, while fusion of local features is not (only SIFT is enough). Fusion of local features, global features, motion features and audio feature achieves the best performance. Post-processing method is very effective to improve overall performance. Mid-level representation shows promising results compared to

using raw feature only, however the performance is still limited. But adding
 335 mid-level feature in fusion schema is not effective. Experimental results on
 MediaEval 2013 VSD benchmark dataset show the validity of the approach and
 its comparable performance to state-of-the-art methods.

References

- [1] C.-H. Demarty, B. Ionescu, Y.-G. Jiang, V. L. Quang, M. Schedl, C. Penet,
 340 Benchmarking violent scenes detection in movies, in: Content-Based Mul-
 timedia Indexing (CBMI), 2014 12th International Workshop on, IEEE,
 2014, pp. 1–6.
- [2] J. Nam, M. Alghoniemy, A. H. Tewfik, Audio-visual content-based violent
 scene characterization, in: Image Processing, 1998. ICIP 98. Proceedings.
 345 1998 International Conference on, Vol. 1, IEEE, 1998, pp. 353–357.
- [3] C. Liang-Hua, H.-W. Hsu, L.-Y. Wang, , C.-W. Su, Violence detection
 in movies, Computer Graphics, Imaging and Visualization (CGIV) (2011)
 119–124.
- [4] C. C., J. Dionisio, M. Echavez, P. C. Naval., Dove: Detection of movie
 350 violence using motion intensity analysis on skin and blood, Workshops and
 Demonstrations - ECCV (2005) 150–156.
- [5] G. Yu, W. Wang, S. Jiang, Q. Huang, W. Gao, Detecting violent scenes in
 movies by auditory and visual cues, Advances in Multimedia Information
 Processing-PCM (2008) 317–326.
- 355 [6] L. Jian, W. Wang, Weakly-supervised violence detection in movies with
 audio and video based co-training, Advances in Multimedia Information
 Processing-PCM (2009) 930–935.
- [7] P. Cdric, C.-H. Demarty, G. Gravier, P. Gros, Technicolor and inria/irisa
 at mediaeval 2011: Learning temporal modality integration with bayesian
 360 networks, MediaEval Multimedia Benchmark Workshop.

- [8] S. Sadanand, J. J. Corso, Action bank: A high-level representation of activity in video, in: Computer Vision and Pattern Recognition (CVPR), 2012 IEEE Conference on, IEEE, 2012, pp. 1234–1241.
- [9] L. Li-Jia, H. Su, L. Fei-Fei, E. P. Xing, Object bank: A high-level image representation for scene classification & semantic feature sparsification, Advances in Neural Information Processing Systems (2010) 1378–1386.
- [10] L. Jingen, B. Kuipers, S. Savarese, Recognizing human actions by attributes, IEEE Conference on Computer Vision and Pattern Recognition (CVPR) (2011) 3337–3344.
- [11] I. Bogdan, J. Schluter, I. Mironica, M. Schedl, A naive mid-level concept-based fusion approach to violence detection in hollywood movies, ACM Conference on International Conference on Multimedia Retrieval (2013) 215–222.
- [12] C. C. Tan, C.-W. Ngo, The vireo team at mediaeval 2013: Violent scenes detection by mid-level concepts learnt from youtube., in: MediaEval, 2013.
- [13] H. Liu, P. Singh, Conceptneta practical commonsense reasoning tool-kit, BT technology journal 22 (4) (2004) 211–226.
- [14] Q. Dai, J. Tu, Z. Shi, Y.-G. Jiang, X. Xue, Fudan at mediaeval 2013: Violent scenes detection using motion features and part-level attributes., in: MediaEval, 2013.
- [15] C. Penet, C.-H. Demarty, G. Gravier, P. Gros, et al., Technicolor/inria team at the mediaeval 2013 violent scenes detection task, MediaEval 2013 Working Notes.
- [16] M. Sjöberg, J. Schlüter, B. Ionescu, M. Schedl, Far at mediaeval 2013 violent scenes detection: Concept-based violent scenes detection in movies., in: MediaEval, 2013.

- [17] N. Derbas, B. Safadi, G. Quénot, et al., Lig at mediaeval 2013 affect task: Use of a generic method and joint audio-visual words., in: MediaEval, Citeseer, 2013.
- 390 [18] C.-C. Chang, C.-J. Lin, LIBSVM: A library for support vector machines, ACM Transactions on Intelligent Systems and Technology 2 (2011) 27:1–27:27, software available at <http://www.csie.ntu.edu.tw/~cjlin/libsvm>.
- 395 [19] V. Lam, D.-D. Le, S.-P. Le, S. Satoh, D. A. Duong, Nii, japan at mediaeval 2012 violent scenes detection affect task., in: MediaEval, Citeseer, 2012.
- [20] H. Wang, A. Kläser, C. Schmid, C.-L. Liu, Action Recognition by Dense Trajectories, in: IEEE Conference on Computer Vision & Pattern Recognition, Colorado Springs, United States, 2011, pp. 3169–3176.
URL <http://hal.inria.fr/inria-00583818/en>